

Flexible Tools for Data Science Education



Damien Eversmann
Chief Architect
Education



Will McGrath
Product Marketing Manager
Cloud Data Services Business Unit

What we will discuss today

- ▶ History of Data Science
- ▶ Problem Statement
- ▶ Elements of a good solution
- ▶ How did Red Hat get involved with Machine Learning?
- ▶ And why Educational institutions?
- ▶ The Future

What is Data Science?



Data science is an interdisciplinary field that uses scientific methods, processes, algorithms and systems to extract knowledge and insights from noisy, structured and unstructured data, and apply knowledge from data across a broad range of application domains. Data science is related to data mining, machine learning and big data.

Where Does It Come From

An Abbreviated History

"Data Analysis"
American statistician, John Tukey, uses the term to define a field close to data science.

Isn't it just Statistics
Statistician Chien-Fu Jeff Wu uses "data science" as a replacement term for "statistics".

It's better than Statistics
Wu again advocates "data science" over "statistics" as the latter just implies "accounting".

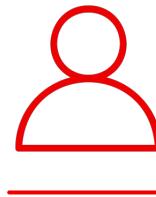
Statistics Applied
Computer Scientist William S. Cleveland proposes "data science" as a term for applied statistics.

Data Scientist
The title begins to be common in the private sector initially at LinkedIn and Facebook.

Fact or Fiction
Some say it is a buzzword like "Big Data", but companies and enterprises everywhere won't make major moves without it.

Intelligent apps are just one of the ways we see Data Science in action

An intelligent app is an app where...



Part of the code was
written by a human



Part of the code is
**a model created from
data and training**

Examples of intelligent applications

- ▶ **Recommendation engines**

Netflix, Amazon, etc..

- ▶ **Virtual assistant**

Siri, Alexa, etc...

- ▶ **Detecting fraudulent activity**

Money laundering, spam, hacking, insurance

- ▶ **Quantifying risks and making smart decisions**

Insurance, loans

- ▶ **Pattern detection**

Images, videos: how many cars, humans, etc. ?

- ▶ **Analyze specialized data**

Seismic data for oil and gas

- ▶ **Teach AI to play video games**

AI opponents

- ▶ **Text analysis**

Summarization, accuracy, offensive, plagiarism detection

- ▶ **Medical**

Tumour detection

- ▶ **Customer retention**

Predict who's about to leave

Poll Question: Do you teach or work in the data science field?



- A. Yes, a significant amount of time
- B. Only a little
- C. Not yet, but looking at doing so in the future
- D. No, mainly interested in learning more

The Problem Defined

Overhead on campus...

"Half my students use the 'laptop-ate-my-homework' excuse to ask for deadline extensions"

Gayatri, Professor

"I waste the first 3 hours of EVERY course helping set up the student's environments."

Sarah, Teaching Assistant

"Every year, I need to do more with less."

Diane, Chief Technology Officer

"Hello IT. Have you tried turning off and on again?"

Roy, Faculty IT

"So of course my laptop decided to blue screen of death an hour before the deadline. I panicked, and now every computer in the house has Pytorch on it."

Ashesh, Undergraduate Student

"I wish I had something else than Excel to teach Linear Regressions"

Pao-lu, Professor

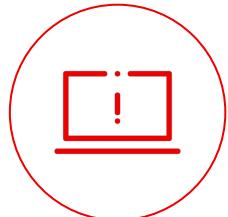
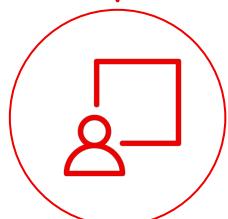
"I've had to rename 'office hours' to 'tech support hours'"

Igor, Adjunct Professor

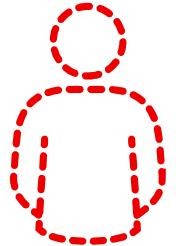
Pain points for educational institutions



- ▶ High variability of needs across faculties, teachers, and students
- ▶ High peaks of activity during live class and the last hour before deadline
- ▶ High variability and low reliability of student-owned devices
- ▶ Different classes can have **different or contradicting software requirements**
- ▶ Waste of time, resources, and talent on zero-value activities
- ▶ Budgets, resources, and skills are split between central IT and faculty IT
- ▶ Rapidly evolving needs based on subjects being taught
 - Data science was not taught 10 years ago
 - MLOps is not taught...yet



What Red Hat tells its commercial companies translates into opportunity for educational institutions



Talent shortage

Lack of key skills makes it difficult to find and retain talent



Lack of self-service access to AI/ML tools & infrastructure

Slows data scientists and developers from doing their job



Complexity to operationalize AI projects

Slow, manual, siloed operations slow AI lifecycle execution

Educating for the future isn't just next gen data scientists



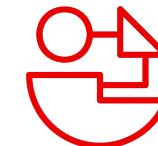
Data
Scientists



Data
Engineers



Developers



Architects

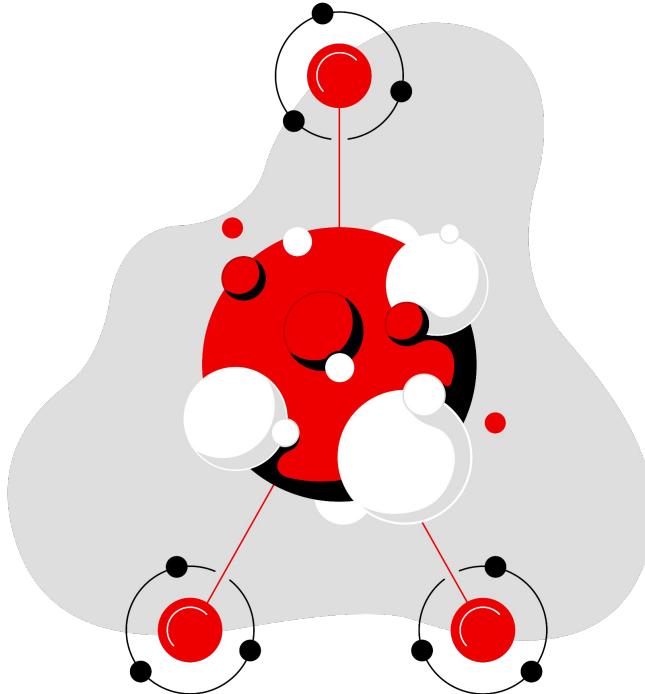


MLOps
Engineers

Elements of a good solution

Open Source

Some of us think this is a no-brainer



- ▶ Open Source is secure and stable
 - ▷ More eyes on the code means fewer flaws and exploits
- ▶ Open Source is nimble and feature rich
 - ▷ Anyone can make changes and participation is encouraged
- ▶ Many/Most of the common Data Science tools are either entirely Open Source or are based in Open Source
- ▶ Most of all, Higher Education is the original Open Source
 - ▷ From Unix in the 70s and GNU in the 80s to Linux and beyond in the 90s, it came from EDU

Flexible and Customizable

Not “One Size Fits None”



- ▶ There are endless tools for Data Science, Artificial Intelligence and Machine Learning.
- ▶ The configurations of those tools increase the permutations exponentially.
- ▶ A good solution should provide an opinionated selection but endless possibilities.



Cloud-Friendly

Available to All

- ▶ Not all schools have multi-million dollar data centers.
 - ▶ They should still be able to teach and sponsor research in Data Science fields
- ▶ Not all students have high-end, multi-GPU personal computers
 - ▶ They should still be able to learn and do Data Science

Rapidly Updating

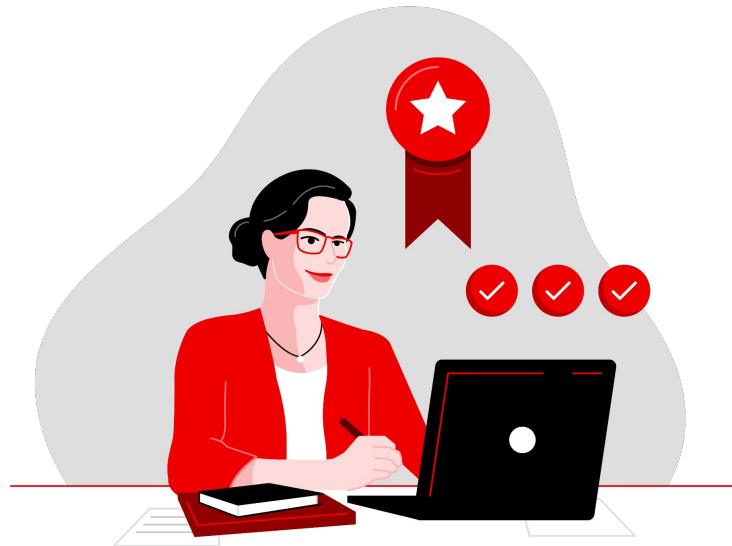
Keeping up with the Joneses



- ▶ The pace of AI/ML research is break neck.
- ▶ The tools that support it appear and change and evolve along with it.
- ▶ A good solution should be able to do the same.

Easy

... as the push of a button



- ▶ These days, everyone is being asked to do more with less.
- ▶ This is especially true with Higher Education...
- ▶ ... And even more so in IT and Data Science.
- ▶ A good solution needs to be simple to install, configure and manage.

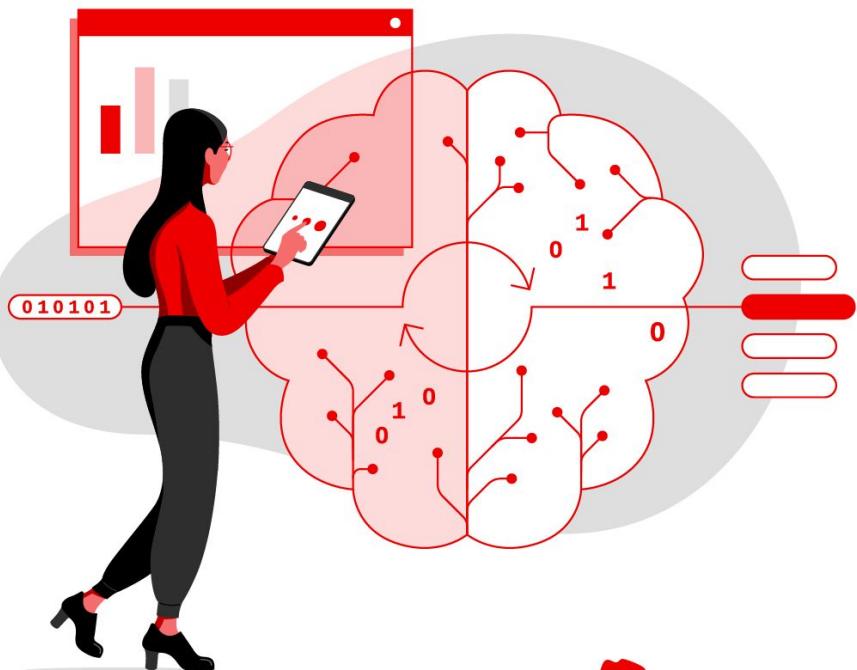
Poll Question: What elements of a data science instructional solution resonate the most with you?



- A. Open source
- B. Flexible and customizable
- C. Cloud friendly
- D. Rapidly updating
- E. Easy
- F. All of above
- G. Other _____

How did Red Hat get involved with Machine Learning?

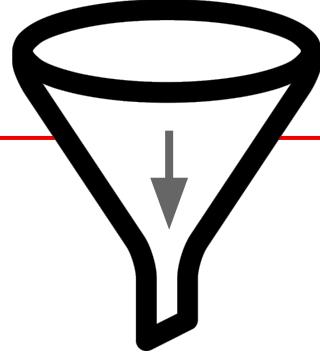
Open Data Hub: the origin story



- ▶ **Began as CI/CD engineering project for build insights**
Terminate batch jobs early
- ▶ **Expanded to an all open source blueprint of AI technologies**
Customer demand from a number of visitors to exec briefing center
- ▶ **Built boutique AI consulting service**
Mainly about helping customers provide AI infra on Kubernetes
- ▶ **Introduced commercial version based on subset of components**
Continued customer ask: can Red Hat support the open source components?

What is Open Data Hub?

100% open source-based ML architecture blueprint built for Kubernetes



Learn more: <https://opendatahub.io>

Based on Open Data Hub and Operate First

Upstream code enhanced with operational excellence

Open Data Hub

Community driven upstream meta-project demonstrating AI/ML platform on Red Hat OpenShift comprised of open source projects

Operate First

Subset Open Data Hub operated at scale for community and university audiences to infuse operational excellence

Red Hat OpenShift Data Science

Subset of Operate First delivered as a cloud service on Red Hat OpenShift
Managed on Amazon Web Services with optional ISV offerings

What is Red Hat OpenShift Data Science

Addressing AI/ML experimentation and integration use cases on a managed platform



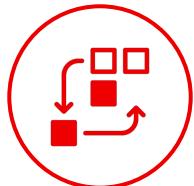
Cloud Service

Available on Red Hat OpenShift Dedicated (AWS) and Red Hat OpenShift Service on AWS



Increased capabilities/collaboration

Combines Red Hat components, open source software, and ISV certified software available on Red Hat Marketplace



Core data science workflow

Provides data scientists and intelligent application developers the ability to build, train, and deploy ML models



Rapid experimentation use cases

Model outputs are hosted on the Red Hat OpenShift managed service or exported for integration into an intelligent application

Dashboard user interface

The screenshot displays the Red Hat OpenShift Data Science dashboard interface. On the left, a sidebar menu includes 'Applications' (with 'Enabled' selected), 'Resources', and 'Settings'. The main content area is divided into two sections: 'Enabled' and 'Explore'.

Enabled: This section shows a summary of enabled applications. It features a card for 'JupyterHub' (Red Hat managed) which is described as a multi-user notebook designed for classrooms and research labs. A 'Launch application' button is present. A blue overlay window titled 'Creating a Jupyter notebook' provides a 'Quick start • 5 minutes' guide.

Explore: This section allows users to add optional applications. It lists several options:

- Anaconda Commercial Edition** (Partner managed): A popular open source package distribution and management experience optimized for commercial use.
- IBM Watson Studio** (Self-managed): A platform for embedding AI and machine learning into business and creating custom models.
- Starburst Galaxy** (Partner managed): A fully managed service for running high-performance queries across various data sources using SQL.
- JupyterHub** (Red Hat managed): A multi-user version of the notebook designed for companies, classrooms, and research labs.
- OpenShift API Management** (Red Hat managed): A service for accelerating time-to-value and reducing the cost of delivering API-first, microservices-based applications.
- OpenShift Streams for Apache Kafka** (Red Hat managed): A service for streaming data that reduces the cost and complexity of delivering real-time applications.
- OpenVINO** (Self-managed): An open source toolkit for optimizing deep learning performance and deploying to Intel hardware.
- Pachyderm** (Self-managed): The data foundation for machine learning, providing industry-leading pipelines, data versioning, and lineage for data science teams to automate the machine learning lifecycle.

Jupyter Spawner - including GPUs

CONFIDENTIAL

Start a notebook server

Select options for your notebook server.

Notebook image

Minimal Python [?](#)

Python v3.8

Standard Data Science [?](#)

Python v3.8

CUDA [?](#)

Python v3.8, CUDA v11.4

Deployment size

Container size

Small

Number of GPUs

0

Number of GPUs

1

0

1

[+ Add more variables](#)

Number of GPUs

0

0

1

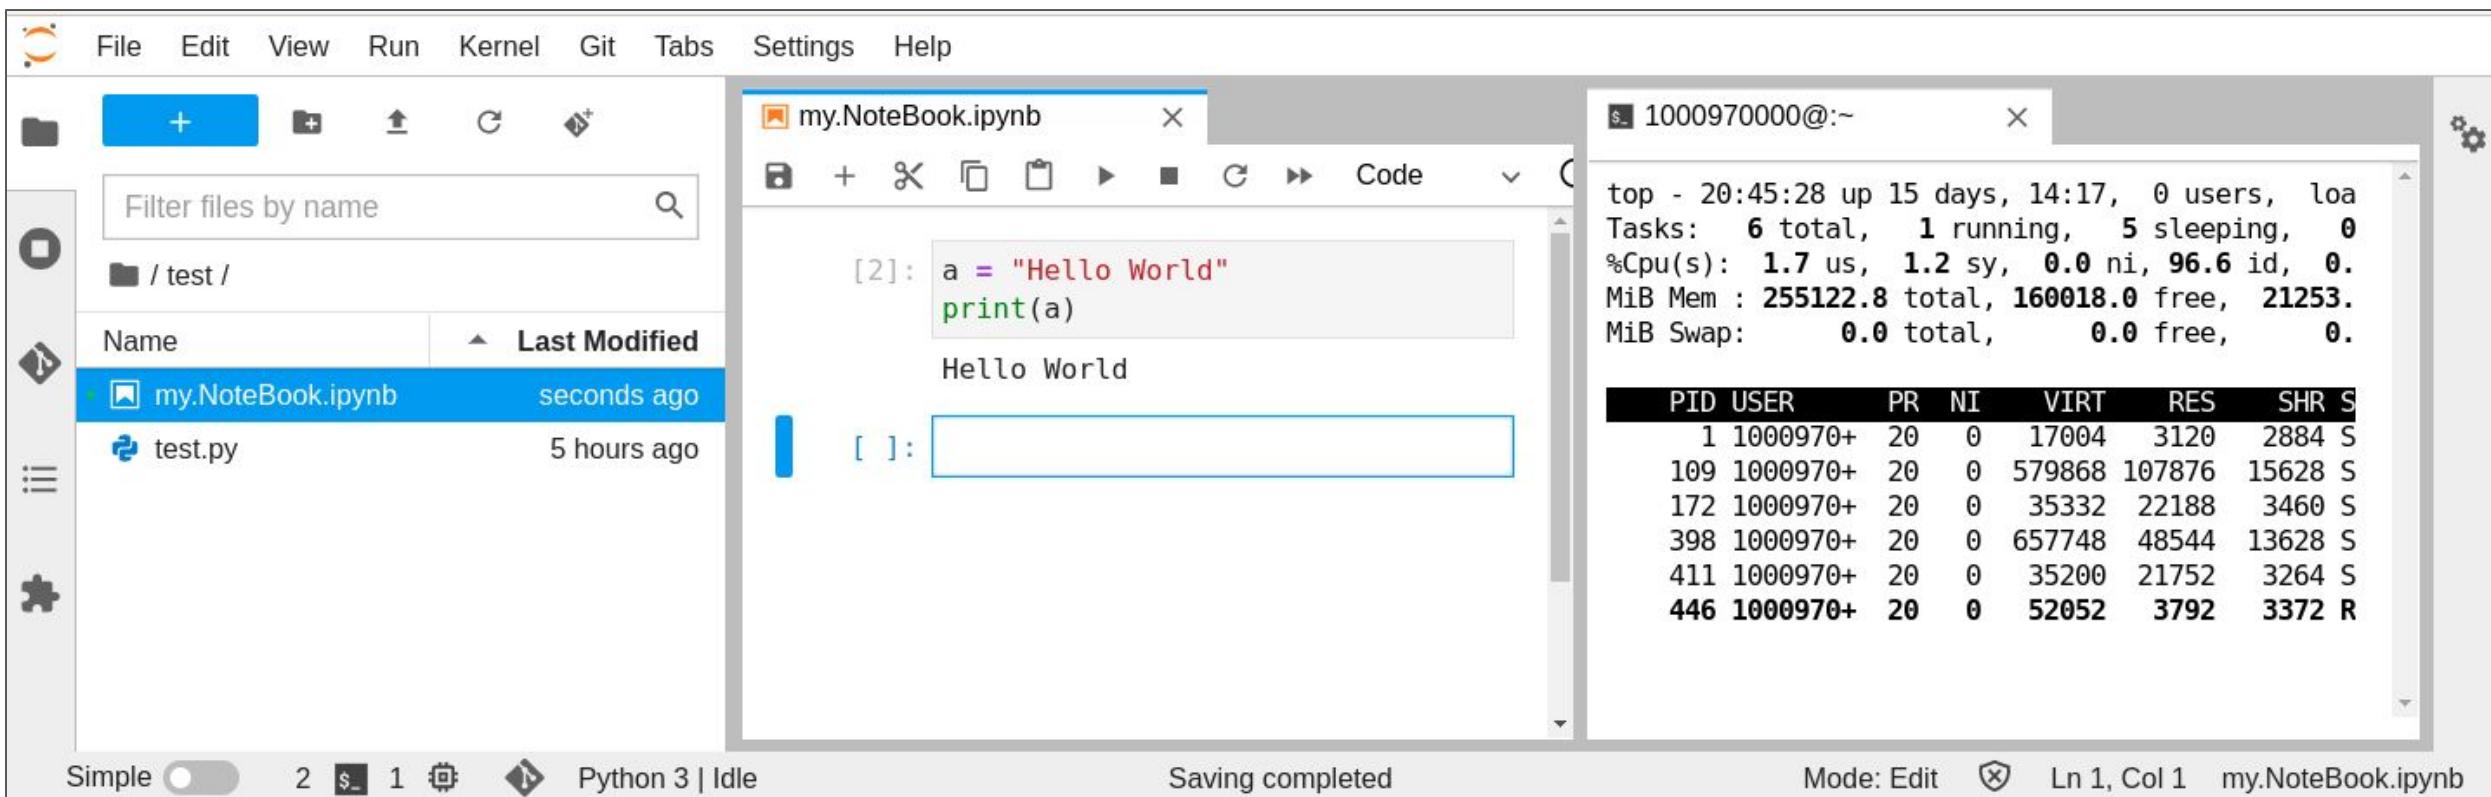
2

3

4

Note: Only way to get GPUs for ROSA in next few months

Standard Jupyter Notebook Interface



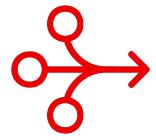
Why Educational Institutions?



- ▶ **Implemented interactive lecture and lab environment** for computer scientists and engineers based on Red Hat OpenShift Data Science
- ▶ **Currently over 300 users** including over 100 concurrent
- ▶ **Integrates with the Boston University online textbook material**, also authored using the Red Hat OpenShift Data Science
- ▶ **Fast time to solution:** cloud services environment enabled BU to configure and deploy in December for classes that started in January
- ▶ **Lowers cost:** auto-scales based on demand; enables bursty interactive use cases at optimized cost

Red Hat OpenShift Data Science

Red Hat® OpenShift® Data Science provides a computing environment for students, faculty, and researchers that is:



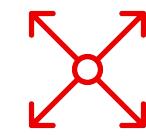
Simple



Managed



Supported



Scalable

Benefits for faculty and research

-  **Spend more time teaching**, less time debugging laptops
-  **Default notebook images**: what the industry uses for data science
-  **(BYO) custom notebook images**: computer science, stats, economics, psychology, etc...
-  **Access** to technology partner software, if required
-  **Consistent**, reliable, and fair environments for all
-  **Not just for class**: research projects, publications, etc...



Benefits to students

 Browser-based environment

 Zero-install, any device

 Environment is **available 24/7**

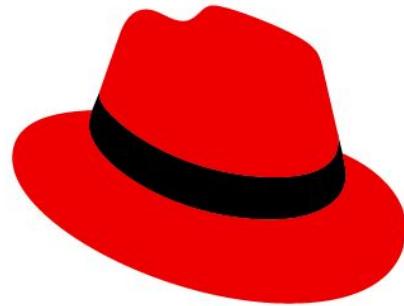
 If it works from your device, it will work from:

- Teacher's device
- Friend's device

 Available during class and for assignments



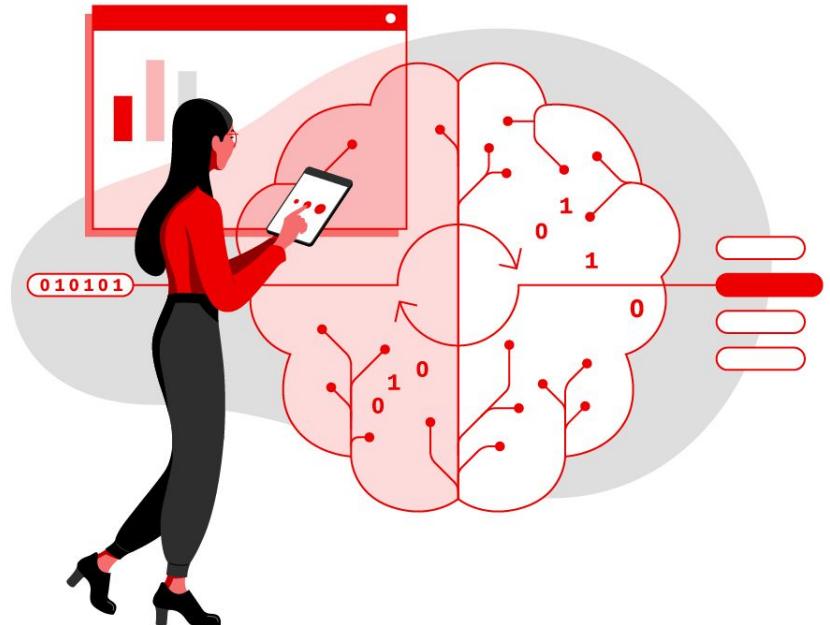
Ready your students ready for their careers



This hat can lead to another

The Future

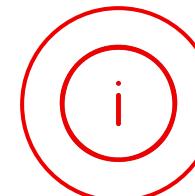
What to expect in coming months



- On prem version - Beta in Sept
- GCP & Azure support
- Better MLOps capabilities through Kubeflow
- Running OpenShift Data Science in same cluster with Open Data Hub components



Red Hat OpenShift Data Science



[Learn more ▶](#)



[Try it ▶](#)

Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



linkedin.com/company/red-hat



youtube.com/user/RedHatVideos



facebook.com/redhatinc



twitter.com/RedHat